

# BACKEND\_STRUCTURE.md — VERDICT

AI-Powered Deposition Coaching & Trial Preparation Platform

Version: 1.0.0 — Hackathon Edition | February 21, 2026

Runtime: Python 3.12 | Framework: FastAPI 0.115.6 | ORM: SQLAlchemy 2.0 + Alembic

## TABLE OF CONTENTS

1. [Architecture Overview](#)
2. [Database Schema](#)
3. [API Endpoints](#)
4. [Authentication & Authorization](#)
5. [Data Validation Rules](#)
6. [Error Handling](#)
7. [Caching Strategy](#)
8. [Rate Limiting](#)
9. [Database Migrations](#)
10. [Backup & Recovery](#)
11. [API Versioning](#)

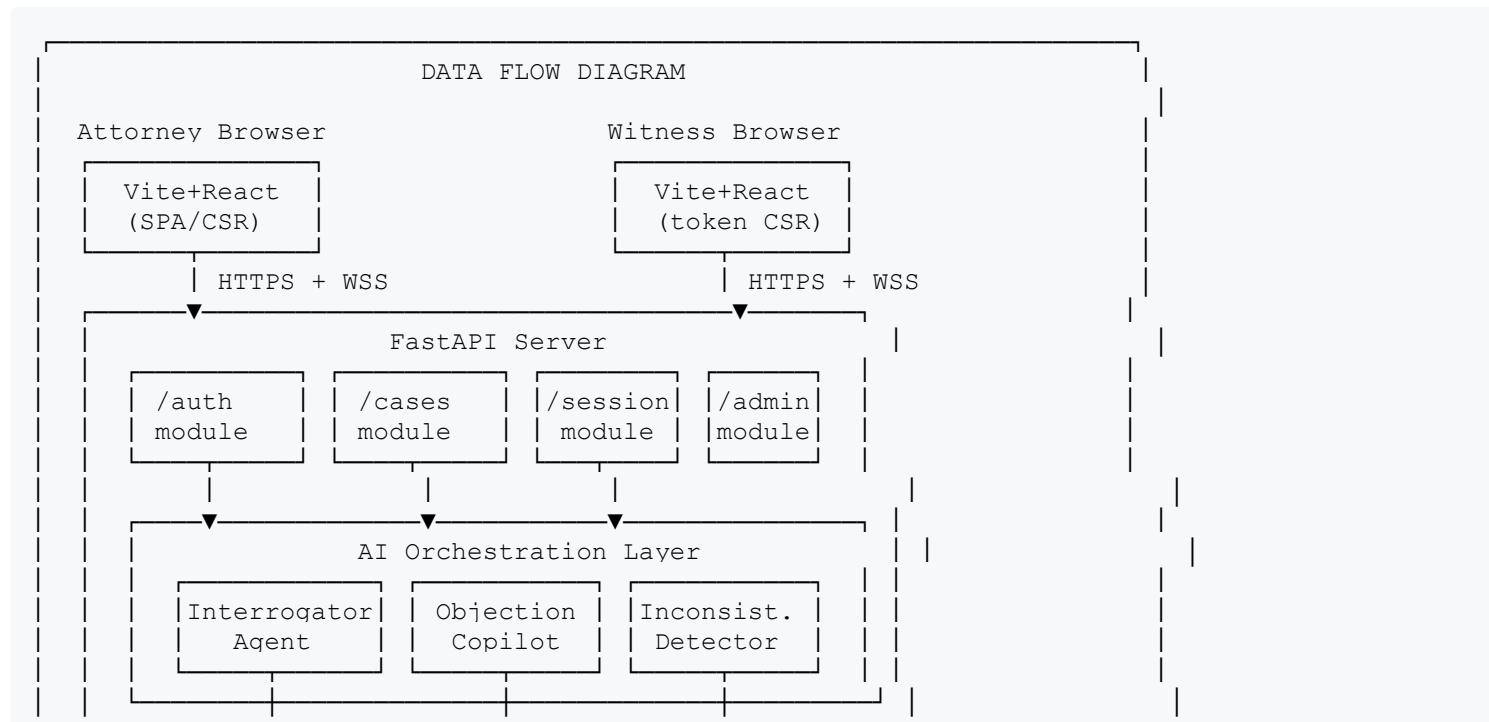
## 1. ARCHITECTURE OVERVIEW

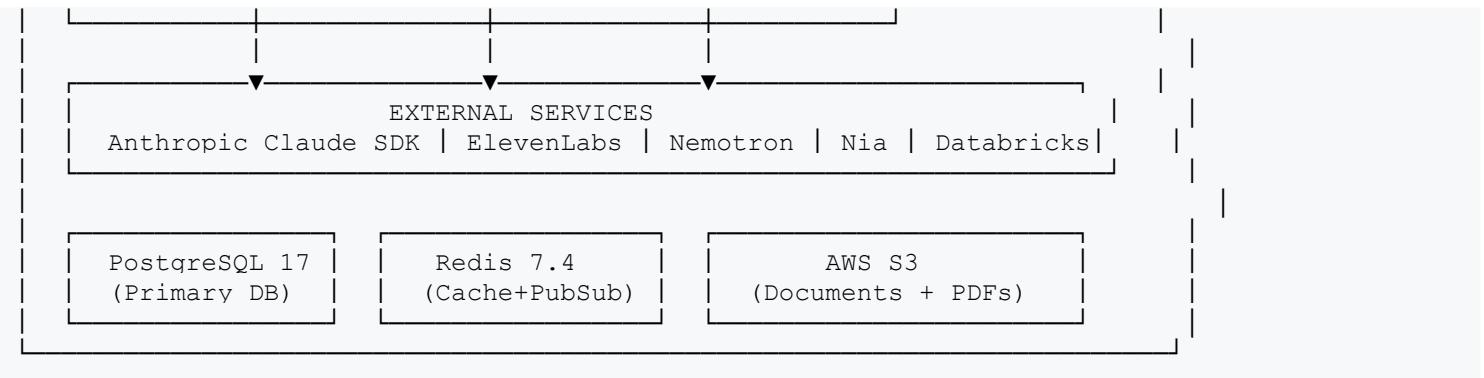
### System Architecture Pattern

#### RESTful API + WebSocket (Socket.io)

VERDICT's backend is a **FastAPI monolith** with **clear module boundaries**, exposing:

- **REST API at `/api/v1/*`** — CRUD operations, document ingestion, brief generation
- **WebSocket namespaces at `/ws/session/:sessionId`** — live session real-time events (alerts, transcript, agent status)
- **AI Agent pipeline** — server-side orchestration of Claude SDK → Nia → Nemotron per session event





## Authentication Strategy

Method	Flow	Storage
<b>JWT (Access Token)</b>	httpOnly cookie, 8h TTL	Client cookie only
<b>JWT (Refresh Token)</b>	httpOnly cookie, 30d TTL	Cookie + Redis allowlist
<b>SAML 2.0</b>	Okta/AzureAD → /auth/saml/callback	Issues JWT on success
<b>Witness Token</b>	nanoid(24) URL param	Redis with 72h TTL
<b>Brief Share Token</b>	nanoid(24) query param	PostgreSQL with expiry timestamp

## Caching Strategy (Summary)

Layer	Store	What	TTL
Session tokens	Redis	Refresh token allowlist	30 days
Witness tokens	Redis	witness:{token} → sessionId	72 hours
Rate limit counters	Redis	Per-IP / per-user endpoint counters	1 minute
Session event buffer	Redis	Live session events (60s auto-flush)	300 seconds
Dashboard stats	Redis	Firm-level aggregated stats	5 minutes
Ingestion status	Redis	ingestion:{documentId} → status + progress	10 minutes

## 2. DATABASE SCHEMA

**Convention:** All tables use PascalCase (as created by Prisma, preserved for compatibility). Column names use camelCase mapped via SQLAlchemy `mapped_column("camelCase", ...)`. All IDs are CUID strings. All timestamps are UTC. All monetary values and scores are stored as integers or floats, never VARCHAR.

### Table: firms

**Purpose:** Stores law firm accounts — the top-level tenant entity. All other data is scoped to a firm.

Column	Type	Constraints	Description
<code>id</code>	VARCHAR(30)	PK, NOT NULL	CUID generated by Prisma
<code>name</code>	VARCHAR(255)	NOT NULL	Law firm display name
<code>slug</code>	VARCHAR(100)	UNIQUE, NOT NULL	URL-safe firm identifier (e.g., "kirkland-ellis")
<code>sso_metadata_url</code>	VARCHAR(2048)	NULL	SAML 2.0 IdP metadata URL for enterprise SSO
<code>sso_metadata_xml</code>	TEXT	NULL	Raw SAML metadata XML (fallback if URL unavailable)
<code>sso_enabled</code>	BOOLEAN	NOT NULL, DEFAULT FALSE	Whether SSO is configured and active
<code>password_login_enabled</code>	BOOLEAN	NOT NULL, DEFAULT TRUE	Whether email/password login is permitted
<code>retention_days</code>	INTEGER	NOT NULL, DEFAULT 90	Days to retain session/case data before deletion
<code>sentinel_enabled</code>	BOOLEAN	NOT NULL, DEFAULT FALSE	Firm-level Behavioral Sentinel (P1.4) toggle
<code>seats</code>	INTEGER	NOT NULL	Maximum number of active users
<code>plan</code>	VARCHAR(50)	NOT NULL, DEFAULT 'trial'	'trial', 'starter', 'professional', 'enterprise'
<code>plan_expires_at</code>	TIMESTAMP	NULL	NULL = no expiry
<code>setup_complete</code>	BOOLEAN	NOT NULL, DEFAULT FALSE	True after onboarding Step 4 completes
<code>created_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Record creation time
<code>updated_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification time

## Indexes:

- PRIMARY KEY (`id`)
- UNIQUE INDEX `firms_slug_unique` (`slug`)
- INDEX `firms_plan_idx` (`plan`)

## Relationships:

- One firm → Many `users` (CASCADE DELETE)
- One firm → Many `cases` (CASCADE DELETE)

## Table: `users`

**Purpose:** Attorneys, associates, paralegals, and admins within a firm.

Column	Type	Constraints	Description
<code>id</code>	VARCHAR (30)	PK, NOT NULL	CUID
<code>firm_id</code>	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Owning firm
<code>email</code>	VARCHAR (255)	UNIQUE, NOT NULL	Login email; RFC 5321 format
<code>name</code>	VARCHAR (255)	NOT NULL	Full display name
<code>role</code>	VARCHAR (30)	NOT NULL	'PARTNER', 'ASSOCIATE', 'PARALEGAL', 'ADMIN'
<code>password_hash</code>	VARCHAR (60)	NULL	bcrypt hash (cost 12); NULL if SSO-only user
<code>password_reset_token</code>	VARCHAR (64)	NULL	Hashed reset token; cleared on use
<code>password_reset_expires_at</code>	TIMESTAMP	NULL	Expiry for password reset token
<code>email_verified</code>	BOOLEAN	NOT NULL, DEFAULT FALSE	True after email verification click
<code>email_verify_token</code>	VARCHAR (64)	NULL	Hashed email verification token; cleared on use
<code>notify_brief_ready</code>	BOOLEAN	NOT NULL, DEFAULT TRUE	Email when coaching brief generated
<code>notify_plateau_alert</code>	BOOLEAN	NOT NULL, DEFAULT TRUE	Email when witness score plateaus
<code>notify_session_reminder</code>	BOOLEAN	NOT NULL, DEFAULT TRUE	Email 24h before session
<code>notify_ingestion_complete</code>	BOOLEAN	NOT NULL, DEFAULT TRUE	Email when document indexing finishes
<code>last_login_at</code>	TIMESTAMP	NULL	Timestamp of most recent successful login
<code>login_attempts</code>	INTEGER	NOT NULL, DEFAULT 0	Failed login counter; reset on success
<code>locked_until</code>	TIMESTAMP	NULL	Account locked until this time (brute-force protection)
<code>is_active</code>	BOOLEAN	NOT NULL, DEFAULT TRUE	FALSE = deactivated (soft delete)
<code>created_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Record creation

updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification
------------	-----------	-------------------------	-------------------

## Indexes:

- PRIMARY KEY (id)
- UNIQUE INDEX users\_email\_unique (email)
- INDEX users\_firm\_id\_idx (firm\_id)
- INDEX users\_firm\_role\_idx (firm\_id, role) — admin panel user listing
- INDEX users\_active\_firm\_idx (firm\_id, is\_active) — seat count queries

## Relationships:

- Many users → One firms (firm\_id)
- One user → Many cases (as owner)
- One user → Many refresh\_tokens

## Table: refresh\_tokens

**Purpose:** Allowlist of active refresh tokens. Enables selective revocation (logout one device) or full revocation (logout all devices).

Column	Type	Constraints	Description
id	VARCHAR (30)	PK, NOT NULL	CUID
user_id	VARCHAR (30)	FK → users.id ON DELETE CASCADE, NOT NULL	Owning user
firm_id	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Firm context (for multi-tenant queries)
token_hash	VARCHAR (64)	UNIQUE, NOT NULL	SHA-256 hash of the raw token
device_hint	VARCHAR (255)	NULL	User-Agent snippet for "active sessions" UI display
ip_address	VARCHAR (45)	NULL	IPv4 or IPv6; stored for security audit
expires_at	TIMESTAMP	NOT NULL	30 days from issuance
revoked_at	TIMESTAMP	NULL	NULL = active; set on logout
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Issuance time
updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

## Indexes:

- PRIMARY KEY (id)
- UNIQUE INDEX refresh\_tokens\_hash\_unique (token\_hash)
- INDEX refresh\_tokens\_user\_id\_idx (user\_id)
- INDEX refresh\_tokens\_expires\_at\_idx (expires\_at) — for nightly expiry cleanup job

## Relationships:

- Many refresh\_tokens → One `users`

## Table: `cases`

**Purpose:** The primary entity — a legal matter for which deposition prep is being conducted.

Column	Type	Constraints	Description
<code>id</code>	VARCHAR (30)	PK, NOT NULL	CUID
<code>firm_id</code>	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Owning firm
<code>owner_id</code>	VARCHAR (30)	FK → users.id ON DELETE RESTRICT, NOT NULL	Attorney who created the case
<code>name</code>	VARCHAR (255)	NOT NULL	e.g., "Chen v. Metropolitan Hospital"
<code>case_type</code>	VARCHAR (50)	NOT NULL	'MEDICAL_MALPRACTICE', 'EMPLOYMENT_DISCRIMINATION', 'COMMERCIAL_DISPUTE', 'CONTRACT_BREACH', 'OTHER'
<code>case_type_custom</code>	VARCHAR (255)	NULL	Free-text for 'OTHER' case type
<code>opposing_firm</code>	VARCHAR (255)	NULL	Opposing counsel firm name
<code>deposition_date</code>	DATE	NULL	Target deposition date for countdown badge
<code>is_archived</code>	BOOLEAN	NOT NULL, DEFAULT FALSE	Soft delete / archive
<code>created_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Record creation
<code>updated_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

## Indexes:

- PRIMARY KEY (`id`)
- INDEX `cases_firm_id_idx` (`firm_id`)
- INDEX `cases_owner_id_idx` (`owner_id`)
- INDEX `cases_firm_archived_idx` (`firm_id, is_archived`) — dashboard listing
- INDEX `cases_deposition_date_idx` (`firm_id, deposition_date`) — sorted by soonest deposition

## Relationships:

- Many cases → One `firms`
- Many cases → One `users` (owner)

- One case → Many documents
- One case → Many witnesses
- One case → Many sessions

## Table: documents

**Purpose:** Case files uploaded by the attorney — PDFs, DOCX, TXT. Tracked through the Nia ingestion pipeline.

Column	Type	Constraints	Description
<code>id</code>	VARCHAR (30)	PK, NOT NULL	CUID
<code>case_id</code>	VARCHAR (30)	FK → cases.id ON DELETE CASCADE, NOT NULL	Owning case
<code>firm_id</code>	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized for row-level security
<code>uploader_id</code>	VARCHAR (30)	FK → users.id ON DELETE SET NULL, NULL	User who uploaded; NULL if user deactivated
<code>filename</code>	VARCHAR (500)	NOT NULL	Original filename as uploaded
<code>s3_key</code>	VARCHAR (1024)	NOT NULL, UNIQUE	S3 object key: <code>{firmId}/{caseId}/{docId}/{filename}</code>
<code>file_size_bytes</code>	BIGINT	NOT NULL	File size in bytes (max 209,715,200 = 200MB)
<code>mime_type</code>	VARCHAR (100)	NOT NULL	'application/pdf', 'application/vnd.openxmlformats-officedocument.wordprocessingml.document', 'text/plain'
<code>doc_type</code>	VARCHAR (50)	NOT NULL	'PRIOR_DEPOSITION', 'MEDICAL_RECORDS', 'FINANCIAL_RECORDS', 'CORRESPONDENCE', 'EXHIBIT', 'OTHER'
<code>page_count</code>	INTEGER	NULL	Set after ingestion succeeds
<code>ingestion_status</code>	VARCHAR (30)	NOT NULL, DEFAULT 'PENDING'	'PENDING', 'UPLOADING', 'INDEXING', 'READY', 'FAILED'
<code>ingestion_started_at</code>	TIMESTAMP	NULL	When Nia pipeline kicked off

ingestion_completed_at	TIMESTAMP	NULL	When status became READY or FAILED
ingestion_error	TEXT	NULL	Error message if status = FAILED
nia_index_id	VARCHAR(255)	NULL	Nia's internal document identifier post-indexing
extracted_facts	JSONB	NULL	{ parties: [], keyDates: [], disputedFacts: [], priorStatements: [] }
facts_confirmed_at	TIMESTAMP	NULL	When attorney confirmed extracted facts
file_hash	VARCHAR(64)	NULL	SHA-256 of file content for duplicate detection
version	INTEGER	NOT NULL, DEFAULT 1	Increments when document is re-uploaded
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Upload initiated
updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last status change

## Indexes:

- PRIMARY KEY (id)
- UNIQUE INDEX documents\_s3\_key\_unique (s3\_key)
- INDEX documents\_case\_id\_idx (case\_id)
- INDEX documents\_firm\_id\_idx (firm\_id)
- INDEX documents\_ingestion\_status\_idx (case\_id, ingestion\_status) — readiness gate queries
- INDEX documents\_file\_hash\_case\_idx (case\_id, file\_hash) — duplicate detection
- GIN INDEX documents\_extracted\_facts\_gin (extracted\_facts) — JSONB search on prior statements

## Relationships:

- Many documents → One cases
- Many documents → One firms

## Table: witnesses

Purpose: Individual deponent records within a case.

Column	Type	Constraints	Description
id	VARCHAR(30)	PK, NOT NULL	CUID
case_id	VARCHAR(30)	FK → cases.id ON DELETE	Owning case

		CASCADE, NOT NULL	
firm_id	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized for isolation
name	VARCHAR (255)	NOT NULL	Witness full name
email	VARCHAR (255)	NOT NULL	Email for practice link delivery
role	VARCHAR (50)	NOT NULL	'DEFENDANT', 'PLAINTIFF', 'EXPERT', 'CORPORATE_REPRESENTATIVE', 'OTHER'
notes	TEXT	NULL	Attorney prep notes (e.g., known weak points)
linked_document_ids	VARCHAR (30) []	NOT NULL, DEFAULT ""	PostgreSQL array of document IDs associated with this witness
session_count	INTEGER	NOT NULL, DEFAULT 0	Denormalized count — incremented on each session creation
latest_score	INTEGER	NULL	Denormalized — updated after each session brief
baseline_score	INTEGER	NULL	Score from Session 1 — never updated
plateau_detected	BOOLEAN	NOT NULL, DEFAULT FALSE	Set to TRUE when no improvement over 3 sessions
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Record creation
updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

## Indexes:

- PRIMARY KEY (id)
- INDEX witnesses\_case\_id\_idx (case\_id)
- INDEX witnesses\_firm\_id\_idx (firm\_id)
- INDEX witnesses\_plateau\_idx (firm\_id, plateau\_detected) — dashboard plateau alerts
- GIN INDEX witnesses\_linked\_doc\_ids\_gin (linked\_document\_ids) — document-witness association queries

## Relationships:

- Many witnesses → One cases
- One witness → Many sessions

---

## Table: sessions

**Purpose:** A single deposition simulation session between an attorney and a witness.

Column	Type	Constraints	Description
id	VARCHAR (30)	PK, NOT NULL	CUID
case_id	VARCHAR (30)	FK → cases.id ON DELETE CASCADE, NOT NULL	Owning case
witness_id	VARCHAR (30)	FK → witnesses.id ON DELETE CASCADE, NOT NULL	Witness being prepared
firm_id	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized
attorney_id	VARCHAR (30)	FK → users.id ON DELETE SET NULL, NULL	Attorney who configured the session
session_number	INTEGER	NOT NULL	1-based sequence per witness (1 = first session)
status	VARCHAR (30)	NOT NULL, DEFAULT 'CONFIGURED'	'CONFIGURED', 'LOBBY', 'ACTIVE', 'PAUSED', 'COMPLETE', 'FAILED'
duration_minutes	INTEGER	NOT NULL	15, 30, 45, or 60
focus_areas	VARCHAR (50) []	NOT NULL, DEFAULT ''	'TIMELINE_CHRONOLOGY', 'FINANCIAL_DETAILS', 'COMMUNICATIONS', 'RELATIONSHIPS', 'ACTIONS_TAKEN', 'PRIOR_STATEMENTS'
aggression_level	VARCHAR (20)	NOT NULL, DEFAULT 'STANDARD'	'STANDARD', 'ELEVATED', 'HIGH_STAKES'
objection_copilot	BOOLEAN	NOT NULL, DEFAULT TRUE	Objection Copilot (P0.2) active
behavioral_sentinel	BOOLEAN	NOT NULL, DEFAULT FALSE	Behavioral Sentinel (P1.4) active
witness_consented	BOOLEAN	NOT NULL, DEFAULT FALSE	Witness granted camera permission for Sentinel
witness_token	VARCHAR (24)	UNIQUE, NULL	nanoid token for witness access link
witness_token_expires_at	TIMESTAMP	NULL	72h from generation

witness_joined_at	TIMESTAMP	NULL	When witness entered the session lobby
started_at	TIMESTAMP	NULL	When attorney clicked "Begin Session"
paused_at	TIMESTAMP	NULL	Most recent pause timestamp
ended_at	TIMESTAMP	NULL	When session completed or was ended early
question_count	INTEGER	NOT NULL, DEFAULT 0	Incrementing counter per question delivered
session_score	INTEGER	NULL	0–100; computed post-session by Review Orchestrator
consistency_rate	FLOAT	NULL	0.0–1.0; proportion of answers with no flags
improvement_delta	INTEGER	NULL	session_score minus witness.baseline_score
transcript_raw	TEXT	NULL	Full speaker-tagged session transcript
nia_session_context_id	VARCHAR(255)	NULL	Nia session context handle for this session's RAG queries
prior_weak_areas	JSONB	NULL	Weak areas from prior sessions injected into Interrogator
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Session record creation
updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

## Indexes:

- PRIMARY KEY (id)
- UNIQUE INDEX sessions\_witness\_token\_unique (witness\_token)
- INDEX sessions\_case\_id\_idx (case\_id)
- INDEX sessions\_witness\_id\_idx (witness\_id)
- INDEX sessions\_firm\_id\_idx (firm\_id)
- INDEX sessions\_attorney\_id\_idx (attorney\_id)
- INDEX sessions\_status\_idx (firm\_id, status) — active session queries
- UNIQUE INDEX sessions\_witness\_number\_unique (witness\_id, session\_number) — enforces sequential numbering

## Relationships:

- Many sessions → One cases
- Many sessions → One witnesses
- Many sessions → One users (attorney)
- One session → Many alerts
- One session → One brief (after completion)
- One session → Many session\_events

## Table: session\_events

**Purpose:** Granular event log for every meaningful action within a live session. Auto-flushed from Redis buffer every 60 seconds. Backs the Databricks Delta Live Tables streaming pipeline.

Column	Type	Constraints	Description
id	VARCHAR(30)	PK, NOT NULL	CUID
session_id	VARCHAR(30)	FK → sessions.id ON DELETE CASCADE, NOT NULL	Owning session
firm_id	VARCHAR(30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized
event_type	VARCHAR(50)	NOT NULL	See event types below
question_number	INTEGER	NULL	Q-number at time of event
speaker	VARCHAR(20)	NULL	'INTERROGATOR', 'WITNESS', 'SYSTEM'
text_content	TEXT	NULL	Question or answer text
audio_latency_ms	INTEGER	NULL	ElevenLabs TTS latency in milliseconds
stt_latency_ms	INTEGER	NULL	ElevenLabs STT latency in milliseconds
emotion_vectors	JSONB	NULL	Behavioral Sentinel: { fear: 0.82, contempt: 0.12, ... }
facs_au_vector	JSONB	NULL	Raw FACS Action Unit values { au4: 0.8, au6: 0.2, ... }
facs_duration_ms	INTEGER	NULL	Duration of detected expression in milliseconds
metadata	JSONB	NULL	Event-type-specific extra data
occurred_at	TIMESTAMP	NOT NULL	Precise event timestamp
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	DB insert time (may be later than occurred_at due to buffer)

### Event Types:

SESSION START	- session began
SESSION PAUSE	- attorney paused
SESSION RESUME	- attorney resumed
SESSION END	- session ended (timer or early)
QUESTION DELIVERED	- Interrogator delivered question (with audio_latency_ms)
ANSWER RECEIVED	- Witness answer transcribed (with stt_latency_ms)
HESITATION DETECTED	- Witness silence > 4 seconds
TOPIC CHANGE	- Interrogator moved to next focus area
OBJECTION FIRED	- Objection Copilot fired alert
INCONSISTENCY FIRED	- Inconsistency Detector fired alert

COMPOSURE ALERT	– Behavioral Sentinel fired (with emotion_vectors, facets_au_vector)
ATTORNEY NOTE	– Attorney timestamped annotation
WITNESS DISCONNECT	– Witness WebSocket dropped
WITNESS RECONNECT	– Witness WebSocket restored
AGENT_DEGRADED	– API fallback activated (e.g., Nemotron → Claude-only)

## Indexes:

- PRIMARY KEY (id)
- INDEX session\_events\_session\_id\_idx (session\_id)
- INDEX session\_events\_firm\_id\_idx (firm\_id)
- INDEX session\_events\_type\_idx (session\_id, event\_type) — event type filtering
- INDEX session\_events\_occurred\_at\_idx (session\_id, occurred\_at) — time-ordered streaming
- GIN INDEX session\_events\_emotion\_vectors\_gin (emotion\_vectors) — Behavioral Sentinel queries

## Relationships:

- Many session\_events → One sessions

## Table: alerts

**Purpose:** Persistent record of every objection, inconsistency, and composure alert — both live-fired and secondary-review-queued.

Column	Type	Constraints	Description
id	VARCHAR(30)	PK, NOT NULL	CUID
session_id	VARCHAR(30)	FK → sessions.id ON DELETE CASCADE, NOT NULL	Owning session
firm_id	VARCHAR(30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized
alert_type	VARCHAR(50)	NOT NULL	'OBJECTION', 'INCONSISTENCY', 'INCONSISTENCY_SECONDARY', 'COMPOSURE'
question_number	INTEGER	NOT NULL	Q-number when alert fired
fired_at	TIMESTAMP	NOT NULL	When the alert was generated
question_text	TEXT	NOT NULL	The question that triggered the analysis
answer_text	TEXT	NULL	Witness answer (for INCONSISTENCY, COMPOSURE)
fre_rule	VARCHAR(50)	NULL	e.g., 'FRE 611(c)' — for OBJECTION alerts
objection_category	VARCHAR(50)	NULL	'LEADING', 'HEARSAY', 'COMPOUND', 'ASSUMES_FACTS', 'SPECULATION'

prior_quote	TEXT	NULL	Exact prior sworn statement for INCONSISTENCY alerts
prior_document_id	VARCHAR(30)	NULL	FK → documents.id (nullable — document may be deleted)
prior_document_page	INTEGER	NULL	Page number in prior document
prior_document_line	INTEGER	NULL	Line number in prior document
contradiction_confidence	FLOAT	NULL	Nemotron score 0.0–1.0
emotion_category	VARCHAR(30)	NULL	'FEAR', 'CONTEMPT', 'DISGUST', 'ANGER', 'SURPRISE' — for COMPOSURE
emotion_duration_ms	INTEGER	NULL	Duration of expression in ms
behavioral_au_vectors	JSONB	NULL	{ au4: 0.8, au20: 0.7, ... }
impeachment_risk	VARCHAR(20)	NULL, DEFAULT 'STANDARD'	'STANDARD', 'HIGH' — HIGH when Sentinel corroborates
is_live_fired	BOOLEAN	NOT NULL, DEFAULT TRUE	FALSE = secondary queue (confidence 0.50–0.74)
attorney_decision	VARCHAR(20)	NULL	'CONFIRMED', 'REJECTED', 'ANNOTATED'
attorney_note	TEXT	NULL	Free-text annotation from attorney
decision_at	TIMESTAMP	NULL	When attorney acted on the alert
in_brief	BOOLEAN	NOT NULL, DEFAULT FALSE	Whether this alert appears in the coaching brief
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Alert record created
updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

## Indexes:

- PRIMARY KEY (id)
- INDEX alerts\_session\_id\_idx (session\_id)
- INDEX alerts\_firm\_id\_idx (firm\_id)
- INDEX alerts\_type\_session\_id\_idx (session\_id, alert\_type) — brief generation queries
- INDEX alerts\_decision\_idx (session\_id, attorney\_decision) — confirmed flags for brief
- INDEX alerts\_persisting\_idx (session\_id, attorney\_decision, alert\_type) — multi-session persisting flags
- INDEX alerts\_fired\_at\_idx (session\_id, fired\_at) — time-ordered alert replay

## Relationships:

- Many alerts → One sessions

## Table: `briefs`

**Purpose:** Post-session coaching brief generated by the Review Orchestrator. One brief per completed session.

Column	Type	Constraints	Description
<code>id</code>	VARCHAR (30)	PK, NOT NULL	CUID
<code>session_id</code>	VARCHAR (30)	FK → sessions.id ON DELETE CASCADE, NOT NULL, UNIQUE	One-to-one with session
<code>firm_id</code>	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized
<code>witness_id</code>	VARCHAR (30)	FK → witnesses.id ON DELETE CASCADE, NOT NULL	Denormalized for witness profile queries
<code>generation_status</code>	VARCHAR (30)	NOT NULL, DEFAULT 'PENDING'	'PENDING', 'GENERATING', 'COMPLETE', 'FAILED'
<code>generation_started_at</code>	TIMESTAMP	NULL	When Review Orchestrator started
<code>generation_completed_at</code>	TIMESTAMP	NULL	When brief became available
<code>generation_error</code>	TEXT	NULL	Error message if FAILED
<code>session_score</code>	INTEGER	NULL	0–100 composite score
<code>consistency_rate</code>	FLOAT	NULL	0.0–1.0
<code>improvement_delta</code>	INTEGER	NULL	vs witness.baseline_score
<code>confirmed_flags</code>	INTEGER	NOT NULL, DEFAULT 0	Count of confirmed inconsistencies
<code>objection_count</code>	INTEGER	NOT NULL, DEFAULT 0	Total objections fired
<code>composure_alert_count</code>	INTEGER	NOT NULL, DEFAULT 0	Total composure alerts
<code>top_recommendations</code>	TEXT []	NOT NULL, DEFAULT "	Array of 3 coaching recommendation strings

narrative_text	TEXT	NULL	Full Claude-generated narrative (attorney-quality prose)
weakness_map_scores	JSONB	NULL	{ timeline: 78, financial: 34, communications: 65, relationships: 80, actions: 71, prior statements: 55, composure: 60 }
elevenlabs_audio_manifest	JSONB	NULL	[ { alertId, s3Key, durationMs } ] — coach voice clips
pdf_s3_key	VARCHAR(1024)	NULL	S3 key for rendered PDF
pdf_generated_at	TIMESTAMP	NULL	When PDF render completed
share_token	VARCHAR(24)	UNIQUE, NULL	nanoid — 7-day expiring share link token
share_token_expires_at	TIMESTAMP	NULL	Expiry for share_token
share_token_issued_at	TIMESTAMP	NULL	When the share token was created
created_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Brief record created
updated_at	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

## Indexes:

- PRIMARY KEY (id)
- UNIQUE INDEX briefs\_session\_id\_unique (session\_id) — one-to-one enforcement
- UNIQUE INDEX briefs\_share\_token\_unique (share\_token)
- INDEX briefs\_firm\_id\_idx (firm\_id)
- INDEX briefs\_witness\_id\_idx (witness\_id) — witness profile brief history
- INDEX briefs\_generation\_status\_idx (generation\_status) — polling for pending briefs
- INDEX briefs\_share\_token\_expiry\_idx (share\_token\_expires\_at) — expired link cleanup

## Relationships:

- One brief ↔ One sessions (1:1)
- Many briefs → One witnesses

## Table: attorney\_annotations

Purpose: Timestamped free-text notes added by attorney during a live session. Separate from alert decisions.

Column	Type	Constraints	Description
id	VARCHAR(30)	PK, NOT NULL	CUID
session_id	VARCHAR(30)	FK → sessions.id ON DELETE CASCADE, NOT NULL	Owning session

<code>firm_id</code>	VARCHAR (30)	FK → firms.id ON DELETE CASCADE, NOT NULL	Denormalized
<code>attorney_id</code>	VARCHAR (30)	FK → users.id ON DELETE SET NULL, NULL	Annotating attorney
<code>question_number</code>	INTEGER	NULL	Q-number at time of annotation
<code>note_text</code>	TEXT	NOT NULL	The annotation content
<code>session_timestamp_ms</code>	INTEGER	NOT NULL	Milliseconds into the session when note was added
<code>created_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Annotation created
<code>updated_at</code>	TIMESTAMP	NOT NULL, DEFAULT NOW()	Last modification

### Indexes:

- PRIMARY KEY (`id`)
- INDEX `annotations_session_id_idx` (`session_id`)

## 3. API ENDPOINTS

**Base URL:** `https://prod.verdict.law/api/v1`

**Content-Type:** `application/json` (all requests and responses)

**Auth header when JWT required:** `Cookie: access_token=Bearer.{jwt}` (`httpOnly`)

**All timestamps:** ISO 8601 UTC (e.g., `"2026-02-21T14:32:00.000Z"`)

## AUTH MODULE

`POST /auth/login`

**Purpose:** Authenticates an attorney or admin with email/password. Issues access + refresh tokens.

**Authentication:** Public

**Request Body:**

```
{
  "email": "sarah.chen@kirklandellis.com",
  "password": "Secure!Pass#2026"
}
```

**Validation Rules:**

- `email`: required, valid RFC 5321 format, max 255 chars
- `password`: required, min 8 chars, max 128 chars

**Response (200):**

```
{
  "success": true,
  "data": {
    "user": {
      "id": "clxyz1234",
      "firmId": "clxyz_firm",
      "email": "sarah.chen@kirklandellis.com",
      "name": "Sarah Chen",
      "role": "PARTNER"
    }
  }
}
```

*Sets two httpOnly cookies:*

- access\_token (8h, Secure, SameSite=Strict)
- refresh\_token (30d, Secure, SameSite=Strict, HttpOnly)

## Errors:

- 400 VALIDATION\_ERROR — Missing or malformed fields
- 401 INVALID\_CREDENTIALS — Email not found OR password mismatch (same message — no enumeration)
- 401 EMAIL\_NOT\_VERIFIED — Account not yet email-verified
- 403 ACCOUNT\_LOCKED — Locked due to too many failed attempts (`locked_until` in future)
- 403 SSO\_REQUIRED — Firm has `sso_enabled=true` and `password_login_enabled=false`
- 403 ACCOUNT\_INACTIVE — `is_active=false`

## Side Effects:

- On success: `last_login_at` updated; `login_attempts` reset to 0; refresh token inserted into `refresh_tokens` table
- On failure: `login_attempts` incremented; if `login_attempts >= 10`: `locked_until` set to `NOW() + 15 minutes`

**POST /auth/logout**

**Purpose:** Revokes the current session's refresh token and clears auth cookies.

**Authentication:** Required

**Request Body:** (none)

**Response (200):**

```
{
  "success": true,
  "message": "Logged out successfully"
}
```

## Side Effects:

- Refresh token row: `revoked_at` set to `NOW()`
- Clears `access_token` and `refresh_token` cookies
- Redis: `DEL session:{userId}:{tokenId}` if cached

**POST /auth/logout-all**

**Purpose:** Revokes ALL refresh tokens for the current user (log out every device).

**Authentication:** Required

**Request Body:** *(none)*

**Response (200):**

```
{  
  "success": true,  
  "message": "All sessions revoked",  
  "sessionsRevoked": 3  
}
```

**Side Effects:**

- All `refresh_tokens` rows for `user_id: revoked_at` set to `NOW()`
- Clears cookies

---

**POST /auth/refresh**

**Purpose:** Issues a new access token using a valid refresh token.

**Authentication:** Public (uses refresh token cookie)

**Request Body:** *(none — reads refresh\_token cookie)*

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "expiresAt": "2026-02-22T22:32:00.000Z"  
  }  
}
```

Sets new `access_token` cookie.

**Errors:**

- 401 TOKEN\_EXPIRED — Refresh token past expiry
- 401 TOKEN\_REVOKED — `revoked_at` is set
- 401 TOKEN\_NOT\_FOUND — No matching hash in `refresh_tokens`
- 403 ACCOUNT\_INACTIVE — User `is_active=false`

---

**GET /auth/saml**

**Purpose:** Initiates SAML 2.0 SSO redirect to the firm's configured IdP.

**Authentication:** Public

**Query Params:**

- `firmSlug` (required) — identifies which IdP to redirect to

**Response:** HTTP 302 redirect to IdP

## Errors:

- 404 FIRM\_NOT\_FOUND — No firm with that slug
  - 400 SSO\_NOT\_CONFIGURED — Firm exists but has no SAML metadata
- 

**POST /auth/saml/callback**

**Purpose:** Handles SAML assertion from IdP. Provisions user if first login. Issues JWT.

**Authentication:** Public (validated via SAML assertion signature)

**Request Body:** SAML XML payload (application/x-www-form-urlencoded from IdP)

**Response:** HTTP 302 redirect to `/dashboard`

## Side Effects:

- If user doesn't exist: provisioned with `role=ASSOCIATE` (admin can upgrade)
  - `last_login_at` updated
  - JWT access + refresh tokens issued
- 

**POST /auth/password/reset-request**

**Purpose:** Sends password reset email. Rate-limited to prevent account enumeration.

**Authentication:** Public

**Request Body:**

```
{  
  "email": "sarah.chen@kirklandellis.com"  
}
```

**Response (200):** Always returns success (prevents email enumeration)

```
{  
  "success": true,  
  "message": "If an account exists, a reset email has been sent."  
}
```

## Side Effects:

- If user found AND firm allows password login: generates reset token, stores SHA-256 hash in `password_reset_token`, sets `password_reset_expires_at` = `NOW() + 1h`
- Sends password reset email via Resend

**Rate Limit:** 5 requests per email per hour

---

**POST /auth/password/reset**

**Purpose:** Completes password reset using the token from email.

**Authentication:** Public

**Request Body:**

```
{  
  "token": "abc123def456...",  
  "newPassword": "NewSecure!Pass#2026"  
}
```

## Validation Rules:

- `token` : required, 64 char hex string
- `newPassword` : min 10 chars, must contain uppercase, lowercase, digit, and special char

## Response (200):

```
{  
  "success": true,  
  "message": "Password updated. Please log in."  
}
```

## Errors:

- 400 `TOKEN_INVALID` — No matching hashed token
- 400 `TOKEN_EXPIRED` — `password_reset_expires_at` in past
- 400 `VALIDATION_ERROR` — Password doesn't meet requirements

## Side Effects:

- New `password_hash` written (bcrypt cost 12)
- `password_reset_token` and `password_reset_expires_at` cleared
- All existing refresh tokens revoked (forces re-login)

## FIRM ONBOARDING MODULE

`POST /onboarding/activate`

**Purpose:** Step 1 of firm setup. Validates the firm invitation token and saves firm configuration.

**Authentication:** Public (requires firm invitation token in request body)

## Request Body:

```
{  
  "invitationToken": "abc123...",  
  "firmName": "Kirkland & Ellis LLP",  
  "adminName": "Sarah Chen",  
  "adminEmail": "sarah.chen@kirklandellis.com",  
  "authMethod": "sso",  
  "ssoMetadataUrl": "https://kirkland.okta.com/app/saml/metadata"  
}
```

## Response (200):

```
{  
  "success": true,  
  "data": {  
    "firmId": "clxyz_firm",  
    "slug": "kirkland-ellis",  
    "setupStep": 2  
  }  
}
```

```
}
```

---

#### POST /onboarding/users

**Purpose:** Step 2 — Bulk provision users via CSV data or individual entries.

**Authentication:** Requires firm admin session (JWT issued after Step 1)

**Request Body:**

```
{
  "users": [
    { "name": "Marcus Webb", "email": "m.webb@kirklandellis.com", "role": "ASSOCIATE" },
    { "name": "Lisa Park", "email": "l.park@kirklandellis.com", "role": "PARALEGAL" }
  ]
}
```

**Response (200):**

```
{
  "success": true,
  "data": {
    "provisioned": 2,
    "skipped": 0,
    "invitationsSent": 2
  }
}
```

---

#### PATCH /onboarding/security

**Purpose:** Step 3 — Saves firm security settings and marks onboarding complete.

**Authentication:** Required (firm admin)

**Request Body:**

```
{
  "retentionDays": 90,
  "sentinelEnabled": false
}
```

**Response (200):**

```
{
  "success": true,
  "data": { "setupComplete": true }
}
```

**Side Effects:**

- `firms.setup_complete` set to TRUE
- Firm record created in Databricks Delta Lake via background job

---

## CASES MODULE

---

## GET /cases

**Purpose:** Lists all cases for the authenticated attorney's firm.

**Authentication:** Required

**Query Params:**

- `archived` (boolean, default: false)
- `sort` ('deposition\_date\_asc', 'created\_at\_desc' — default: 'deposition\_date\_asc')
- `page` (integer, default: 1)
- `limit` (integer, default: 20, max: 50)

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "cases": [  
      {  
        "id": "clxyz_case",  
        "name": "Chen v. Metropolitan Hospital",  
        "caseType": "MEDICAL_MALPRACTICE",  
        "depositionDate": "2026-03-15",  
        "documentsReady": 3,  
        "documentsTotal": 3,  
        "witnessCount": 2,  
        "latestSessionScore": 79,  
        "createdAt": "2026-02-01T09:00:00.000Z"  
      }  
    ],  
    "pagination": {  
      "page": 1,  
      "limit": 20,  
      "total": 7  
    }  
  }  
}
```

**Caching:** Redis key `cases:firm:{firmId}:list` — TTL 5 minutes. Invalidated on: case create, update, archive.

---

## POST /cases

**Purpose:** Creates a new case.

**Authentication:** Required

**Request Body:**

```
{  
  "name": "Chen v. Metropolitan Hospital",  
  "caseType": "MEDICAL_MALPRACTICE",  
  "opposingFirm": "Defense Partners LLP",  
  "depositionDate": "2026-03-15"  
}
```

**Validation Rules:**

- `name` : required, 3–255 chars, sanitized for XSS

- `caseType` : required, must be valid enum value
- `caseTypeCustom` : required when `caseType = "OTHER"` , max 255 chars
- `depositionDate` : optional, must be in future if provided

## Response (201):

```
{
  "success": true,
  "data": {
    "id": "clxyz_case",
    "name": "Chen v. Metropolitan Hospital",
    "caseType": "MEDICAL_MALPRACTICE",
    "depositionDate": "2026-03-15",
    "createdAt": "2026-02-21T14:00:00.000Z"
  }
}
```

## Side Effects:

- Invalidates `cases:firm:{firmId}:list` cache

**GET /cases/:caseId**

**Purpose:** Fetches full case detail including document and witness summary.

**Authentication:** Required (case must belong to authenticated user's firm)

## Response (200):

```
{
  "success": true,
  "data": {
    "id": "clxyz_case",
    "name": "Chen v. Metropolitan Hospital",
    "caseType": "MEDICAL_MALPRACTICE",
    "opposingFirm": "Defense Partners LLP",
    "depositionDate": "2026-03-15",
    "documentsIngestionStatus": "READY",
    "documents": [
      {
        "id": "clxyz_doc",
        "filename": "chen_depo_2024.pdf",
        "docType": "PRIOR_DEPOSITION",
        "ingestionStatus": "READY",
        "pageCount": 247,
        "factsConfirmedAt": "2026-02-10T10:30:00.000Z"
      }
    ],
    "witnesses": [
      {
        "id": "clxyz_witness",
        "name": "Dr. Emily Chen",
        "role": "DEFENDANT",
        "sessionCount": 3,
        "latestScore": 79,
        "plateauDetected": false
      }
    ]
  }
}
```

## Errors:

- 403 FORBIDDEN — Case belongs to different firm
  - 404 NOT\_FOUND — No case with that ID
- 

```
PATCH /cases/:caseId
```

**Purpose:** Updates case name, type, opposing firm, or deposition date.

**Authentication:** Required

**Request Body:**

```
{
  "name": "Chen v. Metropolitan Health System",
  "depositionDate": "2026-03-22"
}
```

**Response (200):** Updated case object (same shape as GET)

**Side Effects:** Cache invalidation for this case and the firm's case list.

---

```
DELETE /cases/:caseId
```

**Purpose:** Soft-archives a case (`is_archived = true`).

**Authentication:** Required (case owner or firm admin)

**Response (200):**

```
{
  "success": true,
  "message": "Case archived"
}
```

## DOCUMENTS MODULE

---

```
POST /cases/:caseId/documents/presign
```

**Purpose:** Generates an S3 pre-signed URL for direct browser-to-S3 upload. Called BEFORE the file upload.

**Authentication:** Required

**Request Body:**

```
{
  "filename": "chen_depo_2024.pdf",
  "mimeType": "application/pdf",
  "fileSizeBytes": 15728640
}
```

**Validation Rules:**

- `filename` : required, max 500 chars

- `mimeType` : must be one of `application/pdf` , `application/vnd.openxmlformats-officedocument.wordprocessingml.document` , `text/plain`
- `fileSizeBytes` : required, max 209,715,200 (200MB)

## Response (200):

```
{
  "success": true,
  "data": {
    "documentId": "clxyz_doc",
    "uploadUrl": "https://s3.amazonaws.com/verdict-docs/...",
    "uploadFields": {},
    "s3Key": "clxyz_firm/clxyz_case/clxyz_doc/chen_depo_2024.pdf",
    "expiresAt": "2026-02-21T15:00:00.000Z"
  }
}
```

## Side Effects:

- Document record created with `ingestion_status = 'UPLOADING'`
- Redis: `SET ingestion:{documentId} → { status: 'UPLOADING', progress: 0 }` TTL 600s

**POST /cases/:caseId/documents/:documentId/confirm-upload**

**Purpose:** Called by frontend after S3 upload completes. Triggers Nia ingestion pipeline.

**Authentication:** Required

**Request Body:** *(none)*

## Response (200):

```
{
  "success": true,
  "data": {
    "documentId": "clxyz_doc",
    "ingestionStatus": "INDEXING"
  }
}
```

## Side Effects:

- Document status updated: `UPLOADING → INDEXING`
- `ingestion_started_at` set to `NOW()`
- Background job queued: pre-screen with `pdf-parse` / `mammoth`, then send to Nia API
- Redis: `ingestion:{documentId}` updated to `{ status: 'INDEXING', progress: 0, pageCount: null }`
- Duplicate hash check: if SHA-256 matches existing doc → return early with `409 DUPLICATE_DOCUMENT`

**GET /cases/:caseId/documents/:documentId/ingestion-status**

**Purpose:** Polls ingestion progress. Frontend polls every 3 seconds while status is INDEXING.

**Authentication:** Required

## Response (200):

```
{
  "success": true,
  "data": {
    "documentId": "clxyz_doc",
    "ingestionStatus": "INDEXING",
    "pagesProcessed": 127,
    "pageCount": 247,
    "progressPercent": 51,
    "eta": "1m 15s"
  }
}
```

**Caching:** Read from Redis `ingestion:{documentId}` first; fall back to PostgreSQL.

**GET /cases/:caseId/documents/facts**

**Purpose:** Returns all extracted facts across all READY documents for attorney review.

**Authentication:** Required

**Response (200):**

```
{
  "success": true,
  "data": {
    "documentsReady": 3,
    "allConfirmed": false,
    "parties": [
      { "id": "p1", "name": "Dr. Emily Chen", "role": "Defendant", "sourceDocId": "clxyz_doc" },
    ],
    "keyDates": [
      { "id": "d1", "date": "2024-03-15", "event": "Initial consultation", "sourceDocId": "clxyz_doc" },
    ],
    "disputedFacts": [
      { "id": "f1", "fact": "Dosage administered was $217", "sourceDocId": "clxyz_doc", "sourcePage": 48 }
    ],
    "priorStatements": [
      { "id": "s1", "quote": "The dosage was exactly $217", "speaker": "Dr. Emily Chen", "sourcePage": 48 }
    ]
  }
}
```

**PATCH /cases/:caseId/documents/facts**

**Purpose:** Saves attorney edits to extracted facts and/or marks a section as confirmed.

**Authentication:** Required

**Request Body:**

```
{
  "updates": [
    { "type": "keyDate", "id": "d1", "changes": { "event": "Initial consultation regarding dosage" } },
    { "type": "priorStatement", "id": "s1", "changes": { "sourcePage": 48 } }
  ],
  "confirmSections": [ "keyDates", "priorStatements" ]
}
```

**Response (200):** Updated facts object (same shape as GET)

---

```
POST /cases/:caseId/documents/:documentId/reingest
```

**Purpose:** Retries ingestion on a FAILED document.

**Authentication:** Required

**Request Body:** (none)

**Response (200):**

```
{  
  "success": true,  
  "data": { "ingestionStatus": "INDEXING" }  
}
```

---

## WITNESSES MODULE

---

```
GET /cases/:caseId/witnesses
```

**Purpose:** Lists all witnesses for a case with session summary.

**Authentication:** Required

**Response (200):**

```
{  
  "success": true,  
  "data": [  
    {  
      "id": "clxyz_witness",  
      "name": "Dr. Emily Chen",  
      "role": "DEFENDANT",  
      "email": "emily.chen@metro-hospital.com",  
      "sessionCount": 3,  
      "latestScore": 79,  
      "baselineScore": 44,  
      "improvementDelta": 35,  
      "plateauDetected": false,  
      "latestSessionId": "clxyz_session3",  
      "latestBriefId": "clxyz_brief3"  
    }  
  ]  
}
```

---

```
POST /cases/:caseId/witnesses
```

**Purpose:** Adds a new witness to a case.

**Authentication:** Required

**Request Body:**

```
{  
  "name": "Dr. Emily Chen",  
  "email": "emily.chen@metro-hospital.com",  
  "role": "DEFENDANT",  
}
```

```
        "notes": "Known weakness: medication dosage timeline",
        "linkedDocumentIds": ["clxyz_doc1", "clxyz_doc2"]
    }
```

## Validation Rules:

- name : required, 2–255 chars, sanitized
- email : required, valid RFC 5321, max 255 chars
- role : required, valid enum
- notes : optional, max 5000 chars
- linkedDocumentIds : must all belong to this case

## Response (201): Created witness object

---

```
GET /cases/:caseId/witnesses/:witnessId
```

**Purpose:** Full witness profile — all sessions, scores, inconsistency log, weakness evolution.

**Authentication:** Required

## Response (200):

```
{
    "success": true,
    "data": {
        "id": "clxyz_witness",
        "name": "Dr. Emily Chen",
        "role": "DEFENDANT",
        "sessionCount": 3,
        "sessions": [
            {
                "id": "clxyz_session1",
                "sessionNumber": 1,
                "score": 44,
                "consistencyRate": 0.61,
                "endedAt": "2026-02-07T18:30:00.000Z",
                "briefId": "clxyz_brief1"
            }
        ],
        "scoreTrend": [44, 61, 79],
        "plateauDetected": false,
        "weaknessMapEvolution": [
            "session1": { "timeline": 40, "financial": 28, "communications": 60, "relationships": 75 },
            "sessionLatest": { "timeline": 72, "financial": 34, "communications": 80, "relationships": 75 }
        ],
        "persistingInconsistencies": [
            {
                "alertId": "clxyz_alert",
                "priorQuote": "The dosage was exactly $217",
                "firstDetectedSession": 1,
                "detectedCount": 2,
                "impeachmentRisk": "HIGH"
            }
        ]
    }
}
```

---

```
PATCH /cases/:caseId/witnesses/:witnessId
```

**Purpose:** Updates witness name, email, role, notes, or linked documents.

## Authentication: Required

---

```
DELETE /cases/:caseId/witnesses/:witnessId
```

**Purpose:** Deletes a witness and all associated sessions, alerts, and briefs.

**Authentication:** Required (case owner or admin only)

**Response (200):**

```
{  
  "success": true,  
  "message": "Witness and all associated session data deleted."  
}
```

---

## SESSIONS MODULE

---

```
POST /cases/:caseId/witnesses/:witnessId/sessions
```

**Purpose:** Creates and configures a new deposition simulation session.

**Authentication:** Required

**Request Body:**

```
{  
  "durationMinutes": 45,  
  "focusAreas": ["TIMELINE_CHRONOLOGY", "FINANCIAL_DETAILS", "PRIOR_STATEMENTS"],  
  "aggressionLevel": "ELEVATED",  
  "objectionCopilot": true,  
  "behavioralSentinel": false  
}
```

---

### Validation Rules:

- `durationMinutes` : required, must be 15, 30, 45, or 60
- `focusAreas` : required, 1–6 items, all valid enum values
- `aggressionLevel` : required, valid enum
- `objectionCopilot` : required, boolean
- `behavioralSentinel` : required, boolean — if true, `firms.sentinel_enabled` must be true

**Response (201):**

```
{  
  "success": true,  
  "data": {  
    "id": "clxyz_session4",  
    "sessionNumber": 4,  
    "status": "CONFIGURED",  
    "witnessToken": null,  
    "priorWeakAreas": {  
      "lowestAxes": ["financial_details"],  
      "persistingInconsistencies": ["dosage timeline"]  
    }  
  }  
}
```

## Side Effects:

- Nia pre-briefing job queued: loads case documents + witness's prior session weak areas into session context

```
POST /sessions/:sessionId/witness-token
```

**Purpose:** Generates a time-limited access token for the witness practice link.

**Authentication:** Required

**Request Body:** *(none)*

**Response (200):**

```
{
  "success": true,
  "data": {
    "witnessToken": "abc123def456ghi789jkl",
    "witnessUrl": "https://verdict.law/witness/session/clxyz_session4?token=abc123...",
    "expiresAt": "2026-02-24T14:00:00.000Z"
  }
}
```

## Side Effects:

- `witness_token` and `witness_token_expires_at` set on session
- Redis: `SET witness:{token} → { sessionId }` EX 259200 (72h)
- Witness invitation email sent via Resend

```
POST /sessions/:sessionId/start
```

**Purpose:** Begins the live session. Transitions status from LOBBY → ACTIVE.

**Authentication:** Required (attorney JWT)

**Request Body:** *(none)*

**Response (200):**

```
{
  "success": true,
  "data": {
    "sessionId": "clxyz_session4",
    "status": "ACTIVE",
    "startedAt": "2026-02-21T15:00:00.000Z"
  }
}
```

## Side Effects:

- `sessions.status → 'ACTIVE'; started_at` set
- WebSocket event broadcast to session room: { `type: 'SESSION_START', sessionId` }
- `session_events` record inserted: type `SESSION_START`
- Interrogator Agent starts question generation pipeline

```
POST /sessions/:sessionId/pause
```

**Purpose:** Pauses session timer and halts Interrogator.

**Authentication:** Required

**Response (200):**

```
{  
  "success": true,  
  "data": { "status": "PAUSED", "pausedAt": "2026-02-21T15:22:00.000Z" }  
}
```

**Side Effects:**

- WebSocket: { type: 'SESSION\_PAUSE' } broadcast to room
- session\_events record: type SESSION\_PAUSE

---

**POST /sessions/:sessionId/resume**

**Purpose:** Resumes paused session.

**Authentication:** Required

**Response (200):** { status: 'ACTIVE' }

---

**POST /sessions/:sessionId/end**

**Purpose:** Ends session and triggers coaching brief generation.

**Authentication:** Required

**Request Body:**

```
{  
  "reason": "TIMER_EXPIRED"  
}
```

**reason** must be "TIMER\_EXPIRED" or "ATTORNEY\_ENDED"

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "status": "COMPLETE",  
    "endedAt": "2026-02-21T15:45:00.000Z",  
    "briefId": "clxyz_brief4",  
    "briefStatus": "GENERATING"  
  }  
}
```

**Side Effects:**

- sessions.status → 'COMPLETE'; ended\_at set
- Transcript finalized and written to sessions.transcript\_raw
- All Redis session event buffer flushed to session\_events table
- Brief record created ( generation\_status = 'GENERATING' )

- Background job started: Review Orchestrator pipeline
- Witness token invalidated in Redis
- WebSocket: `SESSION_END` event + `BRIEF_GENERATING` event
  - witnesses.session\_count incremented; witnesses.latest\_score updated after brief completes

```
GET /sessions/:sessionId
```

**Purpose:** Returns session status, configuration, and live alert summary for the attorney.

**Authentication:** Required

**Response (200):**

```
{
  "success": true,
  "data": {
    "id": "clxyz_session4",
    "status": "ACTIVE",
    "sessionNumber": 4,
    "durationMinutes": 45,
    "startedAt": "2026-02-21T15:00:00.000Z",
    "questionCount": 7,
    "alerts": {
      "objections": 3,
      "inconsistencies": 2,
      "composure": 1
    },
    "agentStatus": {
      "interrogator": "ACTIVE",
      "objectionCopilot": "ACTIVE",
      "inconsistencyDetector": "ACTIVE",
      "behavioralSentinel": "INACTIVE"
    }
  }
}
```

```
GET /witness/sessions/:sessionId (witness-facing)
```

**Purpose:** Validates witness token and returns session metadata for the witness view.

**Authentication:** Witness token (query param `?token=`)

**Response (200):**

```
{
  "success": true,
  "data": {
    "sessionId": "clxyz_session4",
    "status": "LOBBY",
    "caseName": "Confidential Legal Matter",
    "durationMinutes": 45,
    "attorneyName": "Sarah Chen"
  }
}
```

**Errors:**

- 401 TOKEN\_INVALID — No matching token
- 401 TOKEN\_EXPIRED — `witness_token_expires_at` in past

- 410 SESSION\_COMPLETE — Session already finished
- 

## AI AGENTS MODULE

---

```
POST /sessions/:sessionId/agents/question
```

**Purpose:** Generates the next Interrogator question. Called by the session engine after each answer.

**Authentication:** Required (attorney JWT or internal service call)

**Request Body:**

```
{  
  "questionNumber": 8,  
  "priorAnswer": "I always consult before adjusting doses, approximately in the $200 range",  
  "hesitationDetected": false,  
  "behavioralSignal": null,  
  "recentInconsistencyFlag": true,  
  "currentTopic": "FINANCIAL_DETAILS"  
}
```

**Response (200) — Server-Sent Events stream:**

```
data: {"type": "QUESTION START", "questionNumber": 8}  
data: {"type": "QUESTION CHUNK", "text": "Isn't it your sworn testimony"}  
data: {"type": "QUESTION CHUNK", "text": " that the dosage was exactly"}  
data: {"type": "QUESTION CHUNK", "text": " $217?"}  
data: {"type": "QUESTION_END", "fullText": "Isn't it your sworn testimony that the dosage was exa
```

**Side Effects:**

- sessions.question\_count incremented
  - session\_events record: type QUESTION\_DELIVERED , audio\_latency\_ms set after TTS completes
- 

```
POST /sessions/:sessionId/agents/objection
```

**Purpose:** Classifies a question for FRE objection types. Must complete  $\leq 1.5\text{s}$  from question delivery.

**Authentication:** Required

**Request Body:**

```
{  
  "questionNumber": 7,  
  "questionText": "Isn't it true you had forgotten about the dosage by then?",  
  "questionTimestamp": "2026-02-21T15:14:22.000Z"  
}
```

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "isObjectionable": true,  
    "alertId": "clxyz_alert_7",  
  },  
}
```

```

    "category": "LEADING",
    "freRule": "FRE 611(c)",
    "explanation": "Question suggests a specific answer (that the witness had 'forgotten'). Wi",
    "confidence": 0.94,
    "processingMs": 1180
  }
}

```

## Side Effects:

- If `isObjectionable: alerts` record inserted; WebSocket: `objection_alert` event to attorney's socket room

**POST /sessions/:sessionId/agents/inconsistency**

**Purpose:** Checks a witness answer for contradictions against prior sworn statements. Must complete  $\leq 4$ s.

**Authentication:** Required

## Request Body:

```
{
  "questionNumber": 8,
  "questionText": "What was the exact dosage you administered?",
  "answerText": "Approximately $200, in that range",
  "answerTimestamp": "2026-02-21T15:15:03.000Z",
  "behavioralCorroboration": {
    "emotionCategory": "FEAR",
    "durationMs": 1200,
    "auVectors": { "au4": 0.82, "au20": 0.71 }
  }
}
```

## Response (200):

```
{
  "success": true,
  "data": {
    "flagFound": true,
    "alertId": "clxyz_alert_8",
    "isLiveFired": true,
    "contradictionConfidence": 0.91,
    "priorQuote": "The dosage was exactly $217.",
    "priorDocumentId": "clxyz_doc1",
    "priorDocumentPage": 47,
    "priorDocumentLine": 12,
    "impeachmentRisk": "HIGH",
    "behavioralCorroborated": true,
    "processingMs": 3420
  }
}
```

## Decision logic (server-side):

- `confidence >= 0.75 + behavioral corroboration` → `impeachmentRisk = "HIGH"`, live fired
- `confidence >= 0.75 + no corroboration` → `impeachmentRisk = "STANDARD"`, live fired
- `confidence 0.50-0.74` → `isLiveFired = false`, secondary queue
- `confidence < 0.50` → `flagFound = false`, nothing stored

## Side Effects:

- `alerts` record inserted; WebSocket: `inconsistency_alert` event if `isLiveFired`
- 

`POST /sessions/:sessionId/agents/behavioral`

**Purpose:** Receives FACS AU vectors from client-side MediaPipe and classifies emotional state.

**Authentication:** Required

**Request Body:**

```
{  
  "questionNumber": 8,  
  "auVectors": { "au4": 0.82, "au6": 0.11, "au12": 0.05, "au20": 0.71, "au45": 0.02 },  
  "durationMs": 1200,  
  "timestamp": "2026-02-21T15:15:01.000Z"  
}
```

**Validation:** Raw video never accepted — only numeric AU vector JSON.

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "alertFired": true,  
    "alertId": "clxyz_alert_8b",  
    "emotionCategory": "FEAR",  
    "confidence": 0.87,  
    "meetsThreshold": true,  
    "thresholdMs": 800  
  }  
}
```

**Side Effects:**

- If `meetsThreshold`: `alerts` record inserted (type `COMPOSURE`); WebSocket: `composure_alert` event
  - `session_events` record: type `COMPOSURE_ALERT` with `emotion_vectors` and `facs_au_vector`
  - Event streamed to Databricks Delta Live Tables via background flush
- 

## ALERTS MODULE

---

`PATCH /sessions/:sessionId/alerts/:alertId`

**Purpose:** Attorney confirms, rejects, or annotates an alert. Used during session and from the coaching brief.

**Authentication:** Required

**Request Body:**

```
{  
  "decision": "CONFIRMED",  
  "note": "This is the key impeachment moment – drill on this in Session 5"  
}
```

`decision` must be `"CONFIRMED"`, `"REJECTED"`, or `"ANNOTATED"`

## Response (200):

```
{  
  "success": true,  
  "data": {  
    "alertId": "clxyz_alert_8",  
    "decision": "CONFIRMED",  
    "impeachmentRisk": "HIGH",  
    "inBrief": true  
  }  
}
```

## Side Effects:

- alerts.attorney\_decision, attorney\_note, decision\_at updated
- If brief already generated: briefs.confirmed\_flags count updated
- in\_brief flag set based on decision (CONFIRMED → true, REJECTED → false)

## BRIEFS MODULE

GET /briefs/:briefId

**Purpose:** Returns the full coaching brief. Attorney gets full access; share-token access gets read-only subset.

**Authentication:** Required (JWT) OR share token query param ?token=

## Response (200) — attorney view:

```
{  
  "success": true,  
  "data": {  
    "id": "clxyz_brief4",  
    "generationStatus": "COMPLETE",  
    "sessionScore": 79,  
    "consistencyRate": 0.87,  
    "improvementDelta": 35,  
    "confirmedFlags": 2,  
    "objectionCount": 4,  
    "composureAlertCount": 1,  
    "topRecommendations": [  
      "Address the $200 vs $217 dosage discrepancy – prepare a precise, consistent explanation",  
      "Practice a 3-second deliberate pause before answering any leading question.",  
      "Focus next session on Financial Details (weakest axis: 34/100)."  
    ],  
    "narrativeText": "Session 4 showed marked improvement overall...",  
    "weaknessMapScores": {  
      "timeline": 72, "financial": 34, "communications": 80,  
      "relationships": 85, "actions": 78, "prior_statements": 65  
    },  
    "alerts": [  
      {  
        "id": "clxyz_alert_8",  
        "alertType": "INCONSISTENCY",  
        "questionNumber": 8,  
        "questionText": "What was the exact dosage you administered?",  
        "answerText": "Approximately $200, in that range",  
        "priorQuote": "The dosage was exactly $217.",  
        "priorDocumentPage": 47,  
        "priorDocumentLine": 12,  
        "contradictionConfidence": 0.91,  
        "impeachmentRisk": "HIGH",  
      }  
    ]  
  }  
}
```

```

        "behavioralCorroborated": true,
        "attorneyDecision": "CONFIRMED",
        "audioClipS3Key": "firms/clxyz_firm/briefs/clxyz_brief4/alert_8.mp3"
    }
],
"pdfS3Key": "firms/clxyz_firm/briefs/clxyz_brief4/brief.pdf",
"shareToken": "abc123...",
"shareTokenExpiresAt": "2026-02-28T14:00:00.000Z"
}
}

```

**Response (200) — share token (witness/client) view:** Same structure but excludes: `narrativeText` in full (truncated to summary), `alerts[].attorneyDecision`, `alerts[].attorneyNote`, `pdfS3Key` (replaced with public download URL if attorney has shared)

## Errors:

- 401 TOKEN\_INVALID — Invalid share token
- 401 TOKEN\_EXPIRED — Share token past `share_token_expires_at`
- 404 NOT\_FOUND
- 403 BRIEF\_GENERATING — `generation_status` is 'GENERATING' (poll again in 15s)

**POST /briefs/:briefId/pdf**

**Purpose:** Generates or re-generates the PDF for a brief.

**Authentication:** Required

**Request Body:** *(none)*

## Response (202):

```
{
  "success": true,
  "message": "PDF generation started",
  "estimatedSeconds": 30
}
```

## Side Effects:

- Background job: Puppeteer renders brief HTML → PDF → uploads to S3
- On completion: `briefs.pdf_s3_key` and `pdf_generated_at` set; WebSocket: `BRIEF_PDF_READY`

**POST /briefs/:briefId/share**

**Purpose:** Generates a 7-day share token for the brief.

**Authentication:** Required

**Request Body:** *(none)*

## Response (200):

```
{
  "success": true,
  "data": {
    "shareUrl": "https://verdict.law/briefs/clxyz_brief4?token=abc123...",
    "expiresAt": "2026-02-28T14:00:00.000Z"
  }
}
```

```
}
```

## Side Effects:

- `share_token` (nanoid 24) and `share_token_expires_at` set on brief
- Optional: sends share email to witness if email provided

---

```
DELETE /briefs/:briefId/share
```

**Purpose:** Revokes the active share token.

**Authentication:** Required

**Response (200):** { "success": true, "message": "Share link revoked" }

**Side Effects:** `share_token` set to NULL; `share_token_expires_at` cleared.

---

## ADMIN MODULE

---

```
GET /admin/users
```

**Purpose:** Lists all users in the firm for the admin panel.

**Authentication:** Required + Admin role

**Query Params:** `page`, `limit`, `role` filter, `isActive` filter

**Response (200):** Paginated user list with: id, name, email, role, isActive, lastLoginAt, createdAt

---

```
POST /admin/users
```

**Purpose:** Manually adds a single user to the firm.

**Authentication:** Required + Admin role

**Request Body:**

```
{
  "name": "Marcus Webb",
  "email": "m.webb@kirklandellis.com",
  "role": "ASSOCIATE"
}
```

**Side Effects:** User record created; onboarding email sent via Resend.

---

```
PATCH /admin/users/:userId
```

**Purpose:** Updates a user's role or active status.

**Authentication:** Required + Admin role

**Request Body:**

```
{  
  "role": "PARTNER",  
  "isActive": true  
}
```

## Errors:

- 409 SEAT\_LIMIT\_EXCEEDED — Activating a user would exceed firms.seats

---

```
DELETE /admin/users/:userId
```

**Purpose:** Deactivates a user (soft delete — sets `is_active = false`).

**Authentication:** Required + Admin role

**Side Effects:** All refresh tokens for that user revoked; cannot log in again until re-activated.

---

```
GET /admin/stats
```

**Purpose:** Firm-level usage analytics for the admin dashboard.

**Authentication:** Required + Admin role

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "seatsUsed": 18,  
    "seatsTotal": 25,  
    "sessionsThisMonth": 47,  
    "activeCases": 12,  
    "briefs": { "generated": 43, "pdfsDownloaded": 28 },  
    "averageScoreImprovement": 22,  
    "topPerformingAttorney": { "name": "Marcus Webb", "sessionsRun": 14 }  
  }  
}
```

**Caching:** Redis key `admin:stats:firm:{firmId}` — TTL 5 minutes.

---

```
PATCH /admin/settings
```

**Purpose:** Updates firm-level settings (retention, sentinel toggle, SSO config).

**Authentication:** Required + Admin role

**Request Body:**

```
{  
  "retentionDays": 60,  
  "sentinelEnabled": true,  
  "passwordLoginEnabled": false  
}
```

## USER SETTINGS MODULE

---

**GET /users/me**

**Purpose:** Returns the authenticated user's profile and preferences.

**Authentication:** Required

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "id": "clxyz_user",  
    "email": "sarah.chen@kirklandellis.com",  
    "name": "Sarah Chen",  
    "role": "PARTNER",  
    "firm": { "id": "clxyz_firm", "name": "Kirkland & Ellis LLP" },  
    "notifications": {  
      "briefReady": true,  
      "plateauAlert": true,  
      "sessionReminder": true,  
      "ingestionComplete": true  
    },  
    "activeSessions": [  
      { "id": "clxyz_rt1", "deviceHint": "Chrome / Mac", "createdAt": "2026-02-20T08:00:00.000" }  
    ]  
  }  
}
```

**PATCH /users/me**

**Purpose:** Updates name or notification preferences.

**Authentication:** Required

**Request Body:**

```
{  
  "name": "Sarah Chen-Williams",  
  "notifications": {  
    "briefReady": true,  
    "plateauAlert": false  
  }  
}
```

**DELETE /users/me/sessions/:tokenId**

**Purpose:** Revokes a single refresh token (log out one device from active sessions list).

**Authentication:** Required

---

## WEBSOCKET EVENTS

**Connection:** `wss://prod.verdict.law/ws/session/:sessionId`

**Authentication:** JWT must be present in connection handshake cookie. Witness connections include ?token= query param.

### Server → Client Events:

Event Name	Payload	Recipient
session_start	{ sessionId, startedAt }	Both
session_pause	{ pausedAt }	Both
session_resume	{ resumedAt }	Both
session_end	{ endedAt, briefId }	Both
interrogator_question	{ questionNumber, text, audioUrl }	Witness
answer_received	{ questionNumber, transcribedText }	Attorney
objection_alert	{ alertId, questionNumber, category, freRule, explanation }	Attorney only
inconsistency_alert	{ alertId, questionNumber, priorQuote, page, line, confidence, impeachmentRisk }	Attorney only
composure_alert	{ alertId, questionNumber, emotionCategory, durationMs }	Attorney only
witness_connected	{ witnessName }	Attorney only
witness_disconnected	{ questionNumber }	Attorney only
agent_degraded	{ agent, fallbackMode }	Attorney only
brief_generating	{ briefId, estimatedSeconds }	Attorney only
brief_ready	{ briefId }	Attorney only
brief_pdf_ready	{ briefId, pdfUrl }	Attorney only

### Client → Server Events:

Event Name	Payload	Sender
answer_audio	{ audioBlob, questionNumber }	Witness
pause_request	{ }	Both
resume_request	{ }	Attorney
end_request	{ reason }	Attorney

annotation_add	{ questionNumber, text }	Attorney
topic_skip	{ }	Attorney
behavioral_vectors	{ auVectors, durationMs, questionNumber }	Witness (Sentinel)

## 4. AUTHENTICATION & AUTHORIZATION

### JWT Token Structure

#### Access Token Payload:

```
{
  "sub": "clxyz_user",
  "firmId": "clxyz_firm",
  "role": "PARTNER",
  "email": "sarah.chen@kirklandellis.com",
  "iat": 1740139200,
  "exp": 1740168000,
  "jti": "clxyz_jti_access"
}
```

#### Refresh Token Payload:

```
{
  "sub": "clxyz_user",
  "firmId": "clxyz_firm",
  "jti": "clxyz_jti_refresh",
  "iat": 1740139200,
  "exp": 1742731200
}
```

Refresh token only has `sub`, `firmId`, `jti` — no role (fetched fresh on each refresh to pick up role changes).

### Authorization Levels

Level	Routes	Check
Public	/auth/* , /onboarding/* , /witness/sessions/* , /briefs/:id?token=	None
Authenticated	/cases/* , /sessions/* , /agents/* , /alerts/* , /briefs/:id (no token), /users/me	Valid JWT, <code>is_active=true</code>
Firm-scoped	All authenticated routes	<code>resource.firm_id === jwt.firmId</code> — DB-level check
Admin	/admin/*	<code>jwt.role === 'ADMIN'</code>

### FastAPI Auth Middleware

```
# app/middleware/auth.py from fastapi import Depends, HTTPException from fastapi.security import HTTPBearer from jose import JWSError, jwt from sqlalchemy.ext.asyncio import AsyncSession from app.database import get_db from app.models.user import User security = HTTPBearer() async def require_auth(credentials = Depends(security), db: AsyncSession = Depends(get_db)) -> User: try: payload =
```

```

jwt.decode(credentials.credentials, settings.JWT_SECRET, algorithms=["HS256"]) user_id = payload.get("sub")
except JWTError: raise HTTPException(401, detail={"code": "TOKEN_INVALID"}) user = await db.get(User,
user_id) if not user or not user.is_active: raise HTTPException(403, detail={"code": "ACCOUNT_INACTIVE"})
return user

```

## Password Security

Algorithm: bcrypt (passlib[bcrypt] 1.7.4)

Cost factor: 12 (~350ms hash time)

Min rounds: Never go below 10 in any environment

Salt: Auto-generated per hash (not reused)

Password Requirements:

- Minimum 10 characters
- At least one uppercase letter (A-Z)
- At least one lowercase letter (a-z)
- At least one digit (0-9)
- At least one special character (!@#\$%^&\*() \_+=-)
- Maximum 128 characters (bcrypt truncates beyond 72 bytes – enforce max)

Reset Flow:

1. User requests reset → nanoid(32) token generated
2. SHA-256 hash stored in password reset token (never store raw token)
3. Email sent with raw token in URL: /auth/password/reset?token={rawToken}
4. On submit: SHA-256 hash of submitted token compared to stored hash
5. On success: new hash stored, reset token cleared, all refresh tokens revoked
6. Token TTL: 1 hour

## 5. DATA VALIDATION RULES

**Pydantic v2 Schemas (backend `app/schemas/` — Zod used on frontend only)**

```

import { z } from 'zod';

// Email
export const emailSchema = z
  .string()
  .email('Must be a valid email address')
  .max(255, 'Email must be under 255 characters')
  .toLowerCase()
  .trim();

// Password (create/reset)
export const passwordSchema = z
  .string()
  .min(10, 'Password must be at least 10 characters')
  .max(128, 'Password must be under 128 characters')
  .regex(/^[A-Z]/, 'Must contain at least one uppercase letter')
  .regex(/^[a-z]/, 'Must contain at least one lowercase letter')
  .regex(/^[0-9]/, 'Must contain at least one number')
  .regex(/[^!@#$%^&*()_+=-]/, 'Must contain at least one special character');

// Case name
export const caseNameSchema = z
  .string()
  .min(3, 'Case name must be at least 3 characters')
  .max(255, 'Case name must be under 255 characters')
  .transform(s => s.replace(/<[^>]*>/g, '').trim()); // XSS strip

// Witness name
export const witnessNameSchema = z
  .string()
  .min(2, 'Witness name must be at least 2 characters')

```

```

    .max(255, 'Witness name must be under 255 characters')
    .transform(s => s.replace(/<[^>]*>/g, '') .trim()));

// Attorney note / free text (XSS stripped, max length enforced)
export const noteTextSchema = z
  .string()
  .max(5000, 'Note must be under 5000 characters')
  .transform(s => s.replace(/<script\b[^<]*(?:^|<!-->)[^<]*<\\/script>/gi, '') .trim());

// File upload validation (pre-sign step)
export const fileUploadSchema = z.object({
  filename: z.string().min(1).max(500),
  mimeType: z.enum([
    'application/pdf',
    'application/vnd.openxmlformats-officedocument.wordprocessingml.document',
    'text/plain'
  ]),
  fileSizeBytes: z.number().int().positive().max(209_715_200, 'File must be under 200MB'),
});

// Session configuration
export const sessionConfigSchema = z.object({
  durationMinutes: z.union([z.literal(15), z.literal(30), z.literal(45), z.literal(60)]),
  focusAreas: z.array(z.enum([
    'TIMELINE_CHRONOLOGY', 'FINANCIAL_DETAILS', 'COMMUNICATIONS',
    'RELATIONSHIPS', 'ACTIONS_TAKEN', 'PRIOR_STATEMENTS'
  ])).min(1).max(6),
  aggressionLevel: z.enum(['STANDARD', 'ELEVATED', 'HIGH_STAKES']),
  objectionCopilot: z.boolean(),
  behavioralSentinel: z.boolean(),
});

// Behavioral Sentinel AU vectors
export const auVectorSchema = z.object({
  au4: z.number().min(0).max(1),
  au6: z.number().min(0).max(1),
  au12: z.number().min(0).max(1),
  au20: z.number().min(0).max(1),
  au45: z.number().min(0).max(1),
}).strict(); // Reject any unknown keys

```

## Input Length Limits

Field	Min	Max	Notes
Email	5	255	RFC 5321 format validation
Password	10	128	bcrypt 72-byte effective limit
Case name	3	255	XSS stripped
Witness name	2	255	XSS stripped
Attorney notes	0	5,000	XSS stripped
Coaching recommendation	0	1,000	Per recommendation string
Filename	1	500	Sanitized: no path traversal chars
Narrative text (brief)	0	50,000	Claude output
Transcript (raw)	0	500,000	Full session transcript

Prior quote (alert)	1	5,000	Nia-retrieved
AU vector values	0.0	1.0	Float range enforcement

## Content Sanitization

All user-controlled string inputs passed through:

1. **Trim whitespace** — leading/trailing removed
2. **XSS strip** — HTML tags removed via regex before DB write
3. **SQL injection** — SQLAlchemy parameterized queries (no raw SQL with interpolation)
4. **Path traversal** — filenames validated: no .. /, ./, /, \, or null bytes
5. **AI prompt injection** — User-supplied text inserted into Claude prompts inside XML-escaped wrappers:

```
<witness_answer>{ESCAPED_TEXT}</witness_answer>
```

## 6. ERROR HANDLING

### Standard Error Response Format

```
{
  "success": false,
  "error": {
    "code": "INCONSISTENCY_DETECTOR_TIMEOUT",
    "message": "The Inconsistency Detector is taking longer than expected. Falling back to Claude.",
    "statusCode": 503,
    "details": {
      "agent": "inconsistency_detector",
      "fallback": "claude_only",
      "thresholdRaised": true
    },
    "requestId": "req_clxyz_1234",
    "timestamp": "2026-02-21T15:15:03.000Z"
  }
}
```

### Error Code → HTTP Status Mapping

Error Code	HTTP Status	Trigger
VALIDATION_ERROR	400	Zod schema validation failed
MISSING_REQUIRED_FIELD	400	Required field absent
INVALID_ENUM_VALUE	400	Field value not in allowed enum
FILE_TOO_LARGE	400	Upload exceeds 200MB
UNSUPPORTED_FILE_TYPE	400	MIME type not in allowlist
DUPLICATE_DOCUMENT	409	Same file hash already in case
PASSWORD_REQUIREMENTS	400	Password doesn't meet complexity
TOKEN_MISSING	401	No auth token in request
TOKEN_EXPIRED	401	JWT or refresh token past expiry

TOKEN_INVALID	401	Cannot verify signature
TOKEN_REVOKED	401	Refresh token manually revoked
INVALID_CREDENTIALS	401	Login email/password mismatch
EMAIL_NOT_VERIFIED	401	Account not yet verified
ACCOUNT_LOCKED	403	Too many failed login attempts
ACCOUNT_INACTIVE	403	User <code>is_active = false</code>
SSO_REQUIRED	403	Firm requires SSO, password login disabled
FORBIDDEN	403	Resource belongs to different firm
ADMIN_REQUIRED	403	Route requires Admin role
NOT_FOUND	404	Resource doesn't exist
SEAT_LIMIT_EXCEEDED	409	Adding user exceeds firm seat count
SESSION_ALREADY_ACTIVE	409	Witness already has active session
SESSION_COMPLETE	410	Witness token for completed session
BRIEF_GENERATING	202	Brief not yet ready — poll again
RATE_LIMIT_EXCEEDED	429	Too many requests
NEMOTRON_UNAVAILABLE	503	Nemotron API down — fallback active
NIA_UNAVAILABLE	503	Nia API down — fallback active
ELEVENLABS_UNAVAILABLE	503	ElevenLabs API down — text-only mode
INTERNAL_ERROR	500	Unhandled exception

## FastAPI Global Error Handler

```
from fastapi import Request from fastapi.responses import JSONResponse import datetime
@app.exception_handler(Exception) async def global_error_handler(request: Request, exc: Exception): return
JSONResponse( status_code=500, content={ "success": False, "error": { "code": "INTERNAL_ERROR",
"message": str(exc), "statusCode": 500, "timestamp": datetime.datetime.utcnow().isoformat() } } )
```

## 7. CACHING STRATEGY

### Cache Layers & Keys

Cache Key Pattern	Data	TTL	Invalidation Trigger
<code>cases:firm:{firmId}:list</code>	Case list (without full details)	5 min	Case create, update, archive

<code>case:{caseId}:detail</code>	Full case detail + documents + witnesses	2 min	Document upload, witness add/update, session end
<code>ingestion:{documentId}</code>	Ingestion status + progress %	10 min	Nia status webhook; set on READY/FAILED
<code>session:{sessionId}:events</code>	Buffered session events list	5 min	Flushed to DB every 60 seconds
<code>session:{sessionId}:state</code>	Current session state object	5 min	Every state transition
<code>witness:{token}</code>	{ sessionId } for witness token validation	72 hours (259200s)	DEL on session complete or manual revoke
<code>brief:{briefId}:status</code>	Generation status polling	30 sec	Set to COMPLETE on generation finish
<code>admin:stats:firm:{firmId}</code>	Aggregated firm stats	5 min	Any session complete, brief generated
<code>ratelimit:{endpoint}:{identifier}</code>	Request count	60 sec (window-based)	Natural TTL expiry
<code>refresh:{userId}:{jti}</code>	Refresh token validity flag	30 days	DEL on logout

## Session Event Buffer (60-Second Auto-Flush)

```
// Redis list-based event buffer
// Key: session:{sessionId}:events
// Type: Redis List (RPUSH per event, LRANGE + DEL on flush)

const BUFFER_FLUSH_INTERVAL_MS = 60_000;
const EVENT_BUFFER_TTL_S = 300; // 5 min - failsafe TTL

async function bufferSessionEvent(sessionId: string, event: SessionEvent) {
  const key = `session:${sessionId}:events`;
  await redis.rpush(key, JSON.stringify(event));
  await redis.expire(key, EVENT_BUFFER_TTL_S);
}

// Cron-like flush - runs every 60s per active session
async function flushSessionEvents(sessionId: string) {
  const key = `session:${sessionId}:events`;
  const events = await redis.lrange(key, 0, -1);
  if (events.length === 0) return;

  await db.sessionEvent.createMany({
    data: events.map(e => JSON.parse(e))
  });
  await redis.del(key);

  // Also stream to Databricks Delta Live Tables
  await databricks.streamEvents(events.map(e => JSON.parse(e)));
}

// Force-flush on session end (ensures zero data loss)
await flushSessionEvents(sessionId);
```

## Cache Invalidation Triggers

```
// Centralized invalidation - called after any mutation
const cacheInvalidation = {
  onCaseCreate: (firmId: string) => redis.del(`cases:firm:${firmId}:list`),
  onCaseUpdate: (firmId: string, caseId: string) => Promise.all([
    redis.del(`cases:firm:${firmId}:list`),
    redis.del(`case:${caseId}:detail`),
  ]),
  onDocumentIngested: (caseId: string, docId: string) => Promise.all([
    redis.del(`case:${caseId}:detail`),
    redis.del(`ingestion:${docId}`),
  ]),
  onSessionEnd: (firmId: string, caseId: string) => Promise.all([
    redis.del(`case:${caseId}:detail`),
    redis.del(`admin:stats:firm:${firmId}`),
  ]),
  onBriefComplete: (firmId: string) => redis.del(`admin:stats:firm:${firmId}`),
};
```

## 8. RATE LIMITING

### Limits by Endpoint

Route	Limit	Window	Scope	Store
POST /auth/login	10 requests	1 minute	Per IP	Redis
POST /auth/login	5 requests	1 minute	Per email	Redis
POST /auth/refresh	30 requests	1 minute	Per user	Redis
POST /auth/password/reset-request	5 requests	1 hour	Per email	Redis
POST /cases	10 requests	1 minute	Per user	Redis
POST /cases/:id/documents/presign	50 requests	1 hour	Per user	Redis
POST /sessions/:id/agents/question	120 requests	1 minute	Per session	Redis
POST /sessions/:id/agents/objection	120 requests	1 minute	Per session	Redis
POST /sessions/:id/agents/inconsistency	60 requests	1 minute	Per session	Redis
POST /sessions/:id/agents/behavioral	300 requests	1 minute	Per session	Redis
POST /briefs/:id/pdf	10 requests	1 minute	Per user	Redis
GET /api/v1/* (general)	200 requests	1 minute	Per user	Redis

### Rate Limit Response (429)

```
{
  "success": false,
  "error": {
    "code": "RATE_LIMIT_EXCEEDED",
    "message": "Too many requests. Please wait before trying again."}
```

```

    "statusCode": 429,
    "details": {
        "retryAfterSeconds": 42,
        "limit": 10,
        "window": "1 minute"
    },
    "requestId": "req_clxyz_1234",
    "timestamp": "2026-02-21T15:15:03.000Z"
}
}

```

## Response Headers:

```

X-RateLimit-Limit: 10
X-RateLimit-Remaining: 0
X-RateLimit-Reset: 1740142323
Retry-After: 42

```

## FastAPI Rate Limit Configuration

```

from slowapi import Limiter, _rate_limit_exceeded_handler from slowapi.util import get_remote_address from
slowapi.errors import RateLimitExceeded limiter = Limiter(key_func=get_remote_address,
storage_uri=settings.REDIS_URL) app.state.limiter = limiter app.add_exception_handler(RateLimitExceeded,
_rate_limit_exceeded_handler) # Route-specific decorator usage @router.post("/login")
@limiter.limit("10/minute") async def login(request: Request, body: LoginRequest, db: AsyncSession =
Depends(get_db)): ...

```

## 9. DATABASE MIGRATIONS

### Migration Tool: Alembic

#### Migration Naming Convention

```

{NNN} {verb}_{entity}_{field}
Examples:
001 create initial schema
002 add behavioral sentinel columns
003 add session events emotion vectors
004_drop_legacy_alert_status_field (DESTRUCTIVE - two-phase)

```

#### Migration Workflow

```

# Create new migration alembic revision --autogenerate -m "add_behavioral_sentinel_columns" # Apply migrations
alembic upgrade head # Rollback one step alembic downgrade -1 # Check current revision alembic current # View
migration history alembic history

```

#### Migration Process (dev → staging → production)

1. Engineer creates migration on feature branch:  
alembic revision --autogenerate -m "add x"  
→ migration.py generated in alembic/versions/
2. PR opened → CI runs:  
- alembic upgrade head on test DB  
- Full test suite must pass
3. Merge to develop → auto-deploy to staging:

```

- alembic upgrade head
- Staging smoke tests

4. Merge to main → production deploy sequence:
a. alembic upgrade head (runs BEFORE server restart)
b. New FastAPI server starts (new code works with migrated DB)
c. Old server drains connections and shuts down (zero-downtime)

```

## Destructive Change Policy (Two-Phase Deployment)

Problem: Dropping a column crashes old server that still reads it.

Phase 1 (deploy first):

- Add new column OR keep old column but stop writing to it
- Application code reads new column with fallback to old
- Old server remains happy reading old column

Phase 2 (deploy after Phase 1 is stable – next release):

- Drop old column
- Remove fallback read code

Example – renaming alerts.alert type to alerts.category:

```

Phase 1 migration: ADD COLUMN category VARCHAR(50);
                    UPDATE alerts SET category = alert_type;
Phase 1 code: Read category, fallback to alert_type

```

```

Phase 2 migration: DROP COLUMN alert_type;
Phase 2 code: Read category only

```

## Rollback Plan

```

# Alembic supports downgrade steps. # Prevention: ALL migrations must be additive in v1.0 (add columns/tables).
# No DROP COLUMN or DROP TABLE without the two-phase process. # If a bad migration is deployed and must
be undone: 1. Roll back one step: alembic downgrade -1 2. Or roll back to a specific revision: alembic downgrade
001_create_initial_schema 3. Restore from RDS snapshot (nuclear option — last resort): See §10 Backup &
Recovery # Docker image pinning for backend rollback: flyctl deploy --image registry.fly.io/verdict-api:20260221-
stable

```

## Connection Pooling

```

Development: SQLAlchemy create async engine (pool size=5, max_overflow=10)
Production:    asyncpg handles connection pooling natively via SQLAlchemy pool
                  SQLAlchemy create async engine(pool size=10, max_overflow=20)
                  PgBouncer in transaction mode (Supabase pooler)
                  Max PostgreSQL connections: 100

DATABASE URL (production, via Supabase pooler):
postgresql://verdict:pass@aws-pooler.supabase.com:5432/postgres?pgbouncer=true

```

# 10. BACKUP & RECOVERY

## Backup Schedule

Environment	Type	Frequency	Retention	Storage
Hackathon (Supabase)	Managed automated backup	Daily	7 days	Supabase internal

<b>Production (RDS)</b>	Automated RDS snapshot	Daily (2 AM UTC)	35 days	AWS S3 (RDS-managed)
<b>Production (RDS)</b>	Point-in-time recovery (PITR)	Continuous WAL archiving	7 days	AWS S3 (RDS-managed)
<b>Production (S3 docs)</b>	S3 Versioning	On every write	90 days (lifecycle)	S3 versioning
<b>Production (Redis)</b>	ElastiCache snapshot	Daily	7 days	S3
<b>Pre-deploy snapshot</b>	Manual RDS snapshot	Before every production deploy	5 snapshots retained	AWS S3

## Recovery Procedure

### Scenario 1: Accidental Data Deletion (Table-Level)

```
# 1. Identify when deletion occurred from CloudWatch logs / Sentry
# 2. Calculate recovery point: T = now - minutes_since_deletion
# 3. Restore to a read-only RDS instance from PITR: aws rds restore-db-instance-to-point-in-time \
--source-db-instance-identifier verdict-production \
--target-db-instance-identifier verdict-recovery-20260221 \
--restore-time 2026-02-21T14:00:00Z
# 4. Connect to recovery instance and export affected rows:
pg_dump -h verdict-recovery-20260221.us-east-1.rds.amazonaws.com \
-U verdict -d verdict_prod -t alerts --where
"session_id='clxyz'" > alerts_recovery.sql
# 5. Import recovered rows into production: psql -h verdict-production.us-east-1.rds.amazonaws.com \
-U verdict -d verdict_prod < alerts_recovery.sql
# 6. Terminate recovery instance
aws rds delete-db-instance \
--db-instance-identifier verdict-recovery-20260221 \
--skip-final-snapshot
```

### Scenario 2: Full Database Corruption (Catastrophic)

```
# RTO target: < 4 hours | RPO target: < 1 hour (PITR)
# 1. Stop all application traffic (Fly.io scale to 0) flyctl scale count 0 --app verdict-api
# 2. Restore from most recent RDS snapshot (last nightly or pre-deploy): aws rds restore-db-instance-from-db-snapshot \
--db-instance-identifier verdict-production-restored \
--db-snapshot-identifier rds:verdict-production-2026-02-21-02-00
# OR use PITR for more recent recovery point: aws rds restore-db-instance-to-point-in-time \
--source-db-instance-identifier verdict-production \
--target-db-instance-identifier verdict-production-restored \
--restore-time 2026-02-21T14:55:00Z
# 3. Update DATABASE_URL to point to restored instance
# 4. Run alembic upgrade head (verify schema is current)
# 5. Resume traffic flyctl scale count 2 --app verdict-api
```

### Scenario 3: Redis Data Loss (Session Events)

```
# Impact: Up to 60 seconds of session events lost (acceptable per PRD §8.4)
# session_events table has all flushed events (flushed every 60s)
# Restore Redis from ElastiCache snapshot: aws elasticache create-snapshot \
--replication-group-id verdict-redis \
--snapshot-name verdict-redis-recovery
```

## S3 Document Recovery

```
# Documents protected by S3 versioning — restore deleted file: aws s3api restore-object \
--bucket verdict-documents-prod \
--key "clxyz_firm/clxyz_case/clxyz_doc/chen_depo_2024.pdf" \
--version-id "PREVIOUS_VERSION_ID"
```

## 11. API VERSIONING

**Current Version:** v1

## URL Structure

```
https://prod.verdict.law/api/v1/cases  
https://prod.verdict.law/api/v1/sessions/:id  
https://prod.verdict.law/api/v1/briefs/:id
```

## Versioning Strategy

**URL-path versioning** is used (not header-based). The version is the second segment: `/api/v1/`, `/api/v2/` (future).

Rationale: URL-path versioning is unambiguous for enterprise law firm integrations — it doesn't require clients to manage custom headers, making it simpler for firm IT teams and CLI tooling.

## Backwards Compatibility Policy

Change Type	Policy
Add new optional response field	Non-breaking — deploy without version bump
Add new optional request field	Non-breaking — deploy without version bump
Add new endpoint	Non-breaking — deploy without version bump
Change response field name	Breaking → v2 required
Remove response field	Breaking → v2 required
Change field type	Breaking → v2 required
Change HTTP method	Breaking → v2 required
Change auth requirement	Breaking → v2 required

## Deprecation Process

Phase 1 – Announce:

Add Deprecation header to all v1 responses when v2 is released:

Deprecation: true

Sunset: Sat, 22 Feb 2027 00:00:00 GMT

Link: <<https://docs.verdict.law/api/v2/>>; rel="successor-version"

Phase 2 – 6-month window:

Both v1 and v2 run simultaneously

v1 responses include: X-API-Warning: "This version is deprecated..."

Phase 3 – Sunset:

v1 routes return 410 Gone with migration guide URL

Deprecation Notice Channels:

1. HTTP response headers (above)
2. Developer docs changelog at [docs.verdict.law](https://docs.verdict.law)
3. Email to all firm technical contacts on record

## Version Discovery Endpoint

**GET /api/versions**

**Authentication:** Public

**Response (200):**

```
{  
  "success": true,  
  "data": {  
    "current": "v1",  
    "supported": [  
      {  
        "version": "v1",  
        "status": "current",  
        "basePath": "/api/v1",  
        "docsUrl": "https://docs.verdict.law/api/v1",  
        "sunsetAt": null  
      }  
    ]  
  }  
}
```

## QUICK REFERENCE

### Database Tables

Table	Rows (typical firm)	Key Index
firms	1	slug
users	10–25	firm_id, email
refresh_tokens	20–75	user_id, token_hash
cases	15–50	firm_id, deposition_date
documents	50–200	case_id, ingestion_status
witnesses	30–100	case_id, firm_id
sessions	100–500	witness_id, status
session_events	50,000–500,000	session_id, event_type
alerts	1,000–10,000	session_id, alert_type
briefs	100–500	session_id (1:1), witness_id
attorney_annotations	500–5,000	session_id

### AI Agent Latency SLAs

Agent	SLA	Fallback
Interrogator (question gen)	< 2s to audio start	Text-mode if ElevenLabs TTS down
Objection Copilot	≤ 1.5s from question	Claude-only if Nia down

Inconsistency Detector	$\leq 4$ s from answer	Claude-only, threshold 0.85 if Nemotron down
Behavioral Sentinel	$\leq 2$ s FACS → alert	Client-side degradation if camera denied
Review Orchestrator (brief)	$\leq 3$ min post-session	Email when ready if $> 3$ min

## Module File Structure

Current actual structure (`verdict-backend/` — as built):

```

verdict-backend/
└── app/
    ├── init .py
    ├── main.py          # FastAPI entry + CORS + router registration
    ├── config.py        # Pydantic BaseSettings – all env vars
    ├── database.py      # SQLAlchemy async engine + get db()
    ├── redis_client.py  # redis-py async client singleton

    └── models/          # SQLAlchemy ORM (11 models – source of truth)
        ├── enums.py
        ├── firm.py
        ├── user.py
        ├── refresh_token.py
        ├── case.py
        ├── document.py
        ├── witness.py
        ├── session.py
        ├── session_event.py
        ├── alert.py
        ├── brief.py
        └── attorney_annotation.py

    └── schemas/         # Pydantic v2 request/response schemas
        ├── auth.py
        ├── cases.py
        ├── sessions.py
        └── briefs.py

    └── middleware/
        └── auth.py       # require_auth FastAPI Depends()

    └── routers/          # FastAPI APIRouter – all HTTP endpoints
        ├── auth.py
        ├── cases.py
        ├── sessions.py
        └── briefs.py
            # POST /login, POST /logout
            # GET|POST /cases, GET|PATCH|DELETE /cases/:id
            # Session lifecycle + 3 AI agent endpoints
            # GET /briefs/:id, GET /briefs/share/:token

    └── agents/          # AI orchestration layer
        ├── interrogator.py
        ├── objection.py
        ├── detector.py
        └── orchestrator.py
            # Claude streaming question generation
            # FRE classification via Claude
            # Nemotron + Claude fallback contradiction scoring
            # Review Orchestrator – brief generation + TTS

    └── services/         # External API clients
        ├── claude.py
        ├── elevenlabs.py
        ├── nia.py
        └── nemotron.py
            # claude chat() + claude stream()
            # text to speech() + speech to text()
            # index document() + search_index()
            # score_contradiction()

    └── alembic/          # Database migrations
        ├── env.py
        └── versions/
            └── 001_initial_schema.py

    └── scripts/
        └── seed.py        # Dev seed: firm, 3 users, 2 cases, 1 witness

```

```
|  
|   requirements.txt  
|   alembic.ini  
|   run.py          # uvicorn entrypoint  
|   .env.example
```

---

*BACKEND STRUCTURE.md — VERDICT v1.0.0 — Hackathon Edition*

*B2B Deposition Coaching Platform — Team VoiceFlow Intelligence*

*NYU Startup Week Buildathon 2026 | AI Automation — August.law Sponsor Track*